

Doc2vec

Created by: Harsh Bari

From: SVNIT, Gujarat

Mtech Data Science - p23ds004 (2023-25)

Subject: NLP Project

Last Updated: 29/03/2024

```
In [1]: import pandas as pd  
import numpy as np
```


```
In [2]: data = pd.read_csv("resultant_train_dataframe_with_doc2vec.csv")
```

In [3]: data

Out[3]:

	tweets_preprocessed	class	tokens	doc2vec_1	doc2vec_2	doc2vec_3	c
0	aware dirty step get money staylight staywhite...	figurative	['aware', 'dirty', 'step', 'get', 'money', 'st...	0.094989	-0.130547	0.209554	
1	sarcasm people nt understand diy artattack htt...	figurative	['sarcasm', 'people', 'nt', 'understand', 'diy...	0.141715	-0.156211	0.124129	
2	iminworkjeremy medsingle dailymail reader sens...	figurative	['iminworkjeremy', 'medsingle', 'dailymail', '...	0.071912	-0.202637	0.029908	
3	wilw get feeling like game sarcasm	figurative	['wilw', 'get', 'feeling', 'like', 'game', 'sa...	-0.180418	0.049932	-0.014645	
4	teacherarthurg rweingarten probably missed tex...	figurative	['teacherarthurg', 'rweingarten', 'probably', ...	0.315940	-0.020790	-0.059559	
...
81403	photo image via heart http tcoky8nf8z9oi child...	sarcasm	['photo', 'image', 'via', 'heart', 'http', 'tc...	-0.140894	0.102283	0.163148	
81404	never knew better put universe lol maybe date ...	sarcasm	['never', 'knew', 'better', 'put', 'universe',...	-0.286387	0.332654	0.163479	
81405	hey wanted say thanks puberty letting apart it...	sarcasm	['hey', 'wanted', 'say', 'thanks', 'puberty', ...	0.305770	-0.039960	-0.050912	
81406	sure coverage like fox news special hidden har...	sarcasm	['sure', 'coverage', 'like', 'fox', 'news', 's...	0.245714	-0.104332	-0.617117	
81407	skeyno16 u13 wo nt believe see p sarcasm	sarcasm	['skeyno16', 'u13', 'wo', 'nt', 'believe', 'se...	-0.085722	-0.135554	-0.029625	

81408 rows × 153 columns



In [4]: `vec = ['doc2vec_{}'.format(i) for i in range(1, 151)]`
`input_vec = data[vec].values`

In [5]: `# vec`

Split Data

In [6]: `input_vec = np.array(input_vec)`

```
In [7]: def split_data(array_2d, ranges_to_copy):
        copied_ranges = []

        # Loop through each range and copy the corresponding elements
        for start, end in ranges_to_copy:
            copied_range = array_2d[start:end+1] # Adjust end index to include
            copied_ranges.append(copied_range)

        # Concatenate the copied ranges along the first axis to create the final
        copied_array = np.concatenate(copied_ranges, axis=0)

        return copied_array
```

```
In [8]: x_train = split_data(input_vec, [(0, 16989), (21238, 37952), (42132, 57007)])
        x_test = split_data(input_vec, [(16990, 21237), (37953, 42131), (57008, 60700)])
```

```
In [9]: print("x train:", len(x_train))
        print("x test:", len(x_test))
        print("Total:", len(x_train) + len(x_test))
```

```
x train: 65125
x test: 16283
Total: 81408
```

```
In [10]: y_train = np.concatenate((np.zeros(16990), np.ones(31591), np.zeros(16544)))
        y_test = np.concatenate((np.zeros(4248), np.ones(7898), np.zeros(4137)))
```

```
In [11]: print("train:", len(y_train))
        print("test:", len(y_test))
        print("total:", len(y_train) + len(y_test))
```

```
train: 65125
test: 16283
total: 81408
```

Training With Neural Network

```
In [12]: import tensorflow as tf
        from tensorflow import keras
```

```
WARNING:tensorflow:From C:\Users\Harsh Bari\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse_softmax_cross_entropy is deprecated. Please use tf.compat.v1.losses.sparse_softmax_cross_entropy instead.
```

Neural Network for Average Word Embedding

```
In [13]: d2v = keras.Sequential([
    keras.layers.Dense(256, input_shape = (150, ), activation = 'relu'),
    keras.layers.Dense(128, activation = 'relu'),
    keras.layers.Dense(64, activation = 'relu'),
    keras.layers.Dense(32, activation = 'relu'),
    keras.layers.Dense(64, activation=keras.layers.LeakyReLU(alpha=0.1)),
    keras.layers.Dense(32, activation=keras.layers.LeakyReLU(alpha=0.1)),
    keras.layers.Dense(2, activation = 'sigmoid')
])

d2v.compile(optimizer = 'adam',
            loss = 'sparse_categorical_crossentropy',
            metrics = ['accuracy'])
```

WARNING:tensorflow:From C:\Users\Harsh Bari\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.

WARNING:tensorflow:From C:\Users\Harsh Bari\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\optimizers__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.

```
keras.layers.Dense(110, activation = 'relu'), keras.layers.Dense(80,
activation=keras.layers.LeakyReLU(alpha=0.1)),
```

Check Model Summary

In [14]: d2v.summary()

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 256)	38656
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 32)	2080
dense_4 (Dense)	(None, 64)	2112
dense_5 (Dense)	(None, 32)	2080
dense_6 (Dense)	(None, 2)	66
Total params: 86146 (336.51 KB)		
Trainable params: 86146 (336.51 KB)		
Non-trainable params: 0 (0.00 Byte)		

Train Model

```
In [15]: d2v.fit(x_train.astype(np.float32), y_train.astype(np.float32), epochs=22)
```

Epoch 1/22

WARNING:tensorflow:From C:\Users\Harsh Bari\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From C:\Users\Harsh Bari\AppData\Local\Programs\Python\Python310\lib\site-packages\keras\src\engine\base_layer_utils.py:384: The name tf.executing_eagerly_outside_functions is deprecated. Please use tf.compat.v1.executing_eagerly_outside_functions instead.

2036/2036 [=====] - 9s 3ms/step - loss: 0.4561 - accuracy: 0.7652

Epoch 2/22

2036/2036 [=====] - 8s 4ms/step - loss: 0.3947 - accuracy: 0.7964

Epoch 3/22

2036/2036 [=====] - 10s 5ms/step - loss: 0.3678 - accuracy: 0.8101

Epoch 4/22

2036/2036 [=====] - 8s 4ms/step - loss: 0.3453 - accuracy: 0.8208

Epoch 5/22

2036/2036 [=====] - 7s 3ms/step - loss: 0.3273 - accuracy: 0.8287

Epoch 6/22

2036/2036 [=====] - 8s 4ms/step - loss: 0.3120 - accuracy: 0.8355

Epoch 7/22

2036/2036 [=====] - 7s 3ms/step - loss: 0.2977 - accuracy: 0.8415

Epoch 8/22

2036/2036 [=====] - 9s 4ms/step - loss: 0.2861 - accuracy: 0.8466

Epoch 9/22

2036/2036 [=====] - 9s 4ms/step - loss: 0.2744 - accuracy: 0.8505

Epoch 10/22

2036/2036 [=====] - 8s 4ms/step - loss: 0.2652 - accuracy: 0.8546

Epoch 11/22

2036/2036 [=====] - 13s 6ms/step - loss: 0.2586 - accuracy: 0.8576

Epoch 12/22

2036/2036 [=====] - 11s 5ms/step - loss: 0.2511 - accuracy: 0.8598

Epoch 13/22

2036/2036 [=====] - 11s 5ms/step - loss: 0.2418 - accuracy: 0.8637

Epoch 14/22

2036/2036 [=====] - 8s 4ms/step - loss: 0.2364 - accuracy: 0.8661

Epoch 15/22

2036/2036 [=====] - 9s 5ms/step - loss: 0.2321 - accuracy: 0.8679

Epoch 16/22

2036/2036 [=====] - 10s 5ms/step - loss: 0.2266 - accuracy: 0.8706

Epoch 17/22

2036/2036 [=====] - 7s 4ms/step - loss: 0.2217 - accuracy: 0.8726

```

Epoch 18/22
2036/2036 [=====] - 7s 3ms/step - loss: 0.2176 -
accuracy: 0.8757
Epoch 19/22
2036/2036 [=====] - 8s 4ms/step - loss: 0.2119 -
accuracy: 0.8773
Epoch 20/22
2036/2036 [=====] - 7s 3ms/step - loss: 0.2078 -
accuracy: 0.8801
Epoch 21/22
2036/2036 [=====] - 7s 4ms/step - loss: 0.2043 -
accuracy: 0.8816
Epoch 22/22
2036/2036 [=====] - 10s 5ms/step - loss: 0.2009 -
accuracy: 0.8839

```

Out[15]: <keras.src.callbacks.History at 0x24968989a80>

Training Accuracy

In [16]: `d2v.evaluate(x_train.astype(np.float32), y_train.astype(np.float32))`

```

2036/2036 [=====] - 6s 3ms/step - loss: 0.1951 -
accuracy: 0.8917

```

Out[16]: [0.19513297080993652, 0.8917466402053833]

Testing Accuracy

In [17]: `prediction = d2v.predict(x_test.astype(np.float32))`

```

509/509 [=====] - 1s 2ms/step

```

In [18]: `prediction = np.argmax(prediction, axis = 1)`

In [19]: `from sklearn.metrics import classification_report, confusion_matrix, accuracy_score`


```
In [20]: print(classification_report(y_test.astype(np.float32), prediction))
print()
print("Confusion Matrix: \n", confusion_matrix(y_test.astype(np.float32), p
print("\nAccuracy: \n", accuracy_score(y_test.astype(np.float32), predictio
```

	precision	recall	f1-score	support
0.0	0.84	0.76	0.80	8385
1.0	0.77	0.85	0.81	7898
accuracy			0.81	16283
macro avg	0.81	0.81	0.81	16283
weighted avg	0.81	0.81	0.81	16283

Confusion Matrix:

```
[[6414 1971]
 [1198 6700]]
```

Accuracy:

```
0.8053798440090892
```